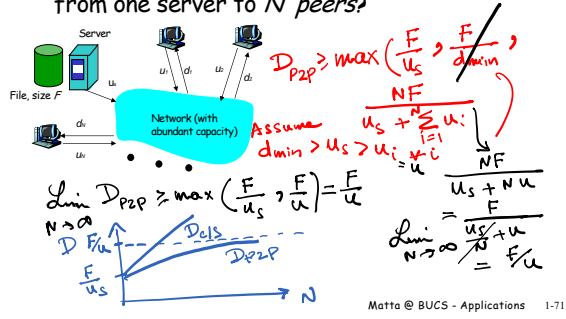


File Distribution: Server-Client vs P2P

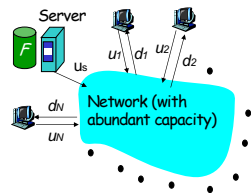
Question: How much time to distribute file from one server to N peers?



71

File distribution time: client-server

- server sequentially sends N copies:
 - NF/u_s time
- client i takes F/d_i time to download



Time to distribute F to N clients using client/server approach

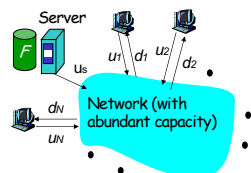
increases linearly in N (for large N)

Matta © BUCS - Applications 1-72

72

File distribution time: P2P

- server must send one copy: F/u_s time
- client i takes F/d_i time to download
- NF bits must be downloaded (aggregate)
 - fastest possible upload rate: $u_s + \sum u_i$



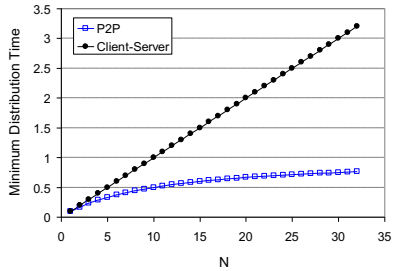
$$d_{P2P} \geq \max\left\{\frac{F}{u_s}, \frac{F}{\min(d_i)}, \frac{NF}{u_s + \sum u_i}\right\}$$

Matta © BUCS - Applications 1-73

73

Server-client vs. P2P: example

Client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{\min} \geq u_s$



Matta @ BUCS - Applications 1-74

74

File distribution: BitTorrent

P2P file distribution

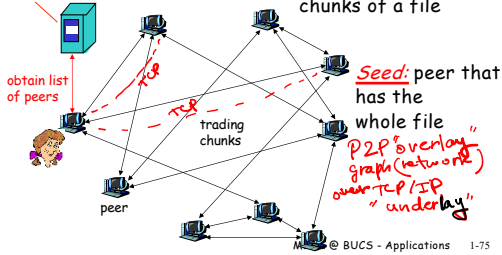
tracker: tracks peers participating in torrent

obtain list of peers

file name \rightarrow list of IP addresses of existing peers

torrent: group of peers exchanging chunks of a file

Seed: peer that has the whole file
P2P "overlay" graph (network) over TCP/IP "underlay"

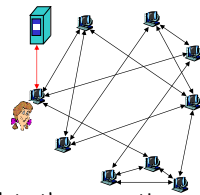


Matta @ BUCS - Applications 1-75

75

BitTorrent (1)

- file divided into 256KB *chunks*
- peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- while downloading, peer uploads chunks to other peers
- peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain



Matta @ BUCS - Applications 1-76

76

BitTorrent (2)

① chunk/piece selection

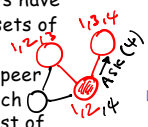
Pulling Chunks

- at any given time, different peers have different subsets of file chunks
- periodically, a peer (Alice) asks each neighbor for list of chunks that they have
- Alice sends requests for her missing chunks
 - rarest first policy

② peer selection

Sending Chunks: tit-for-tat

- Alice sends chunks to four neighbors currently sending her chunks *at the highest rate*
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - newly chosen peer may join top 4
 - "optimistically unchoke"

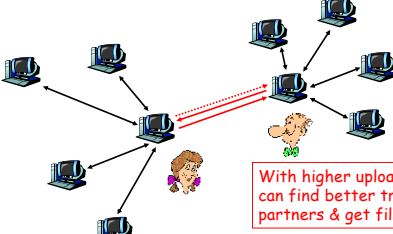


Matta © BUCS - Applications 1-77

77

BitTorrent: Tit-for-tat

- Alice "optimistically unchokes" Bob
- Alice becomes one of Bob's top-four providers; Bob reciprocates
- Bob becomes one of Alice's top-four providers

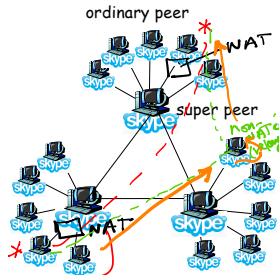


Matta © BUCS - Applications 1-78

78

P2P Case study: Skype

- proprietary application-layer protocol (inferred via reverse engineering)
- Hierarchical overlay of Skype peers
- Index maps usernames to IP addresses; distributed over super peers

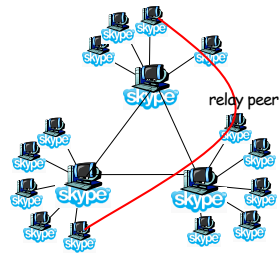


Matta © BUCS - Applications 1-79

79

Peers as relays

- ❑ Problem when both Alice and Bob are behind "NATs"
 - NAT prevents an outside peer from initiating a call to insider peer
- ❑ Solution:
 - Using Alice's and Bob's super peers, (non-NATed) Relay is chosen
 - Each peer initiates session with relay
 - Peers can now communicate through NATs via relay



Matta @ BUCS - Applications 1-80

80

Chapter 2: Summary

our study of network apps now complete!

- ❑ application architectures
 - client-server
 - P2P
 - hybrid
- ❑ application service requirements:
 - reliability, throughput, delay
- ❑ Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- ❑ socket programming
- ❑ specific protocols:
 - HTTP
 - SMTP, POP, IMAP
 - DNS
 - P2P: BitTorrent, Skype, ...

Matta @ BUCS - Applications 1-81

81

Chapter 2: Summary

Most importantly: learned about protocols

Important themes:

- ❑ typical request/reply message exchange:
 - ▢ client requests info or service
 - ▢ server responds with data, status code
- ❑ message formats:
 - ▢ headers: fields giving info about data
 - ▢ data: info being communicated
- ❑ persistent vs. non-persistent transport connections
- ❑ stateless vs. stateful
- ❑ caching
- ❑ reliable vs. unreliable msg transfer
- ❑ centralized vs. distributed
- ❑ Overlay vs. underlay

Matta @ BUCS - Applications 1-82

82