

CL NL Service Datagrams

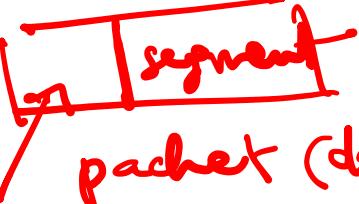
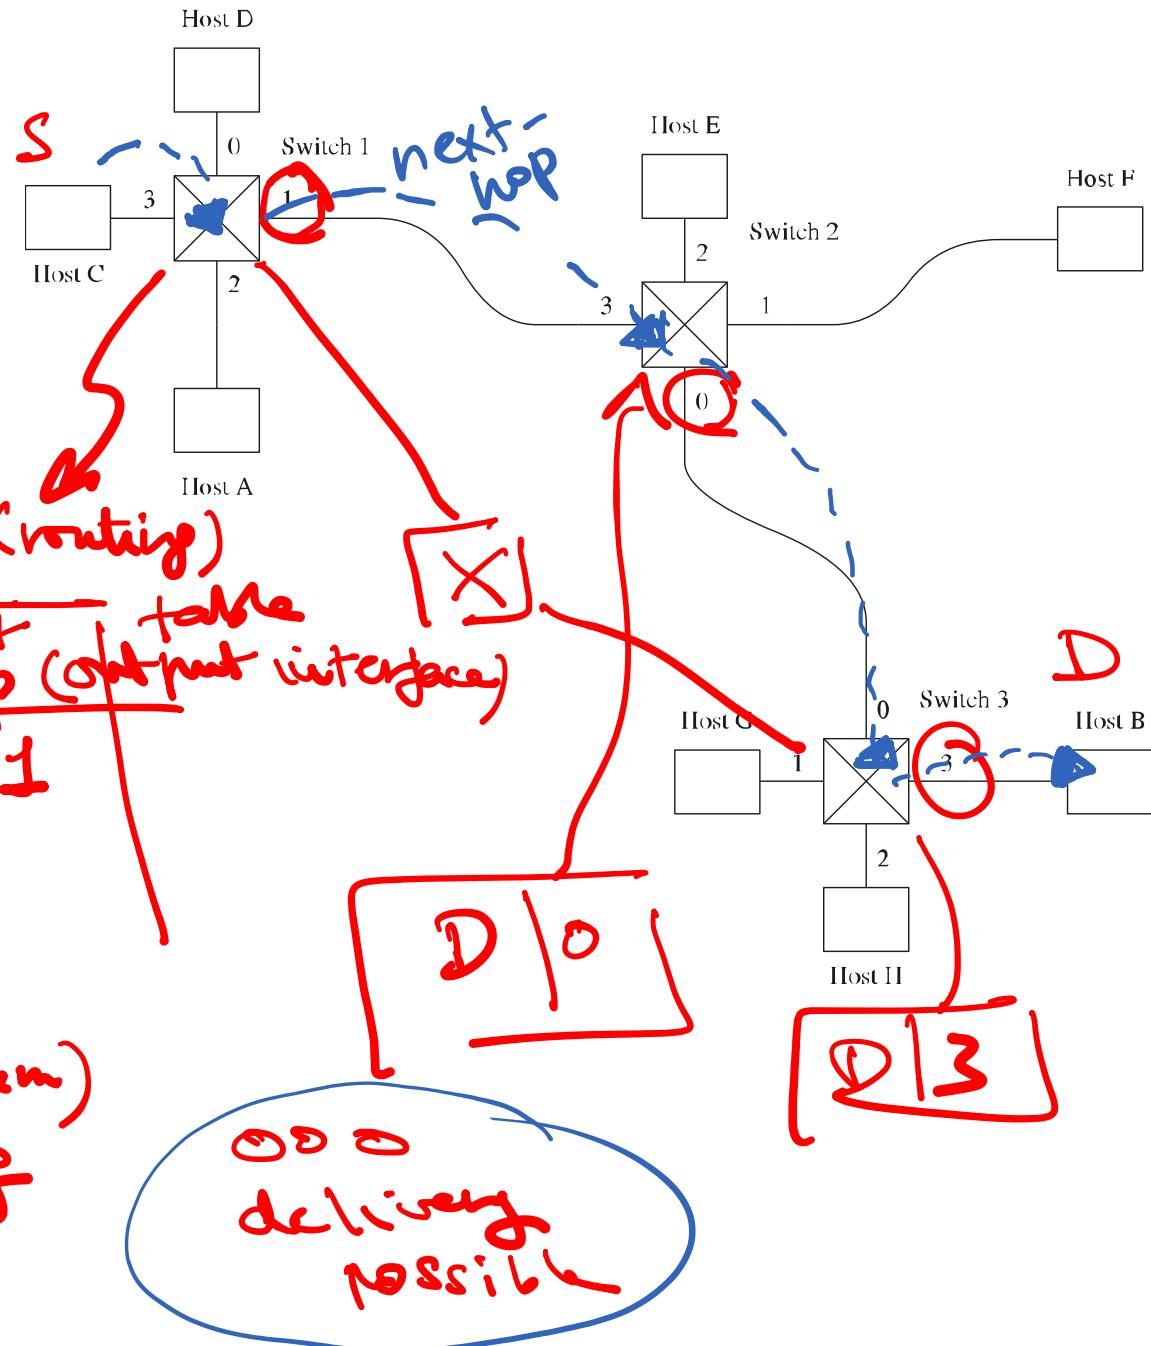
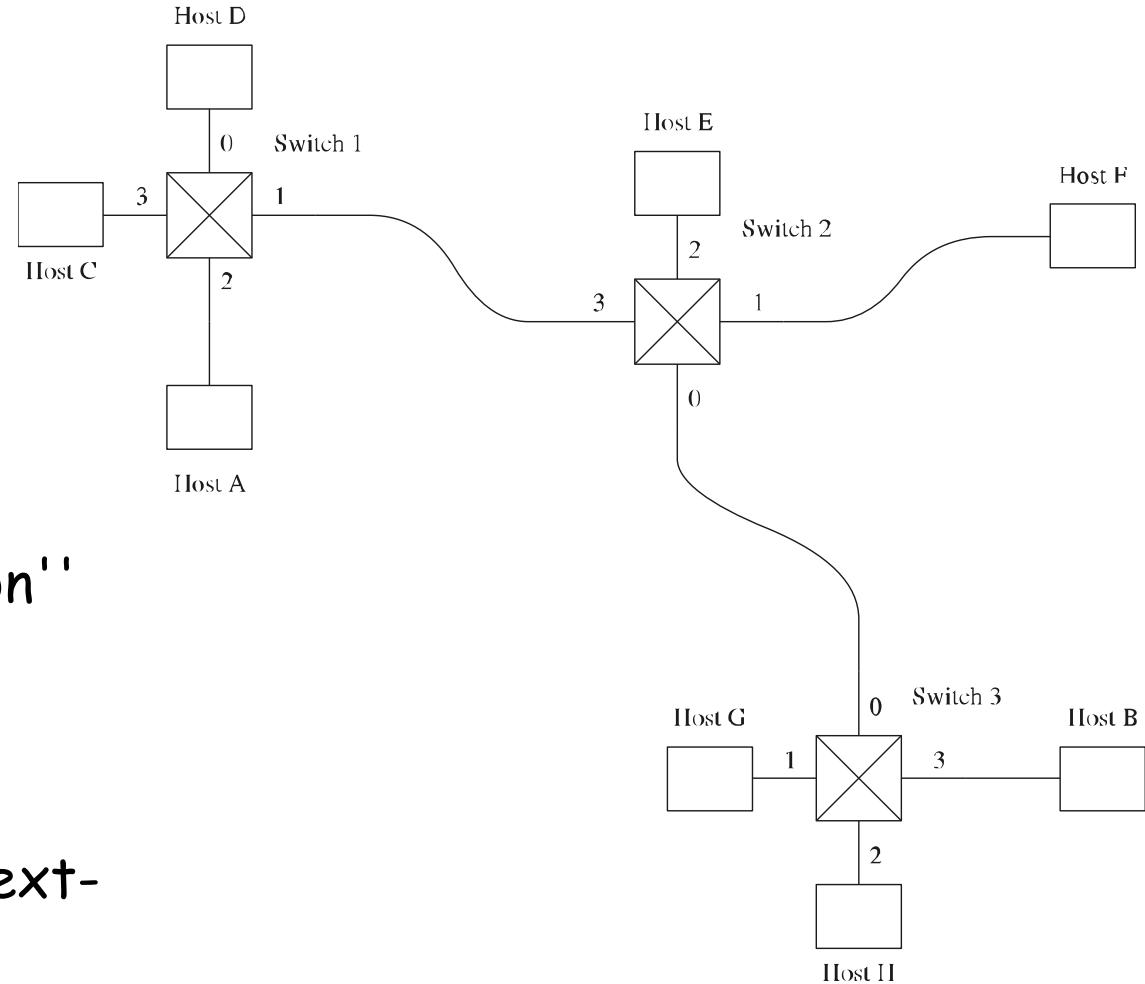

 packet (datagram)
 dest address
 forwarding (routing)

 table
 Dest address | next-hop (output interface)
 D | 1
 next-hop (datagram) matching



Datagrams

- ❑ No connection setup phase
- ❑ Each packet forwarded independently based on destination address
- ❑ Each packet in ``connection'' may follow different path
- ❑ Analogy: postal system
- ❑ Sometimes called ***connectionless*** model (or next-hop routing)
- ❑ Each switch maintains a forwarding (routing) table:
{destination address, output port (or next-hop)}



Virtual Circuits

dest address
setup packet
next-hop routed

VC table

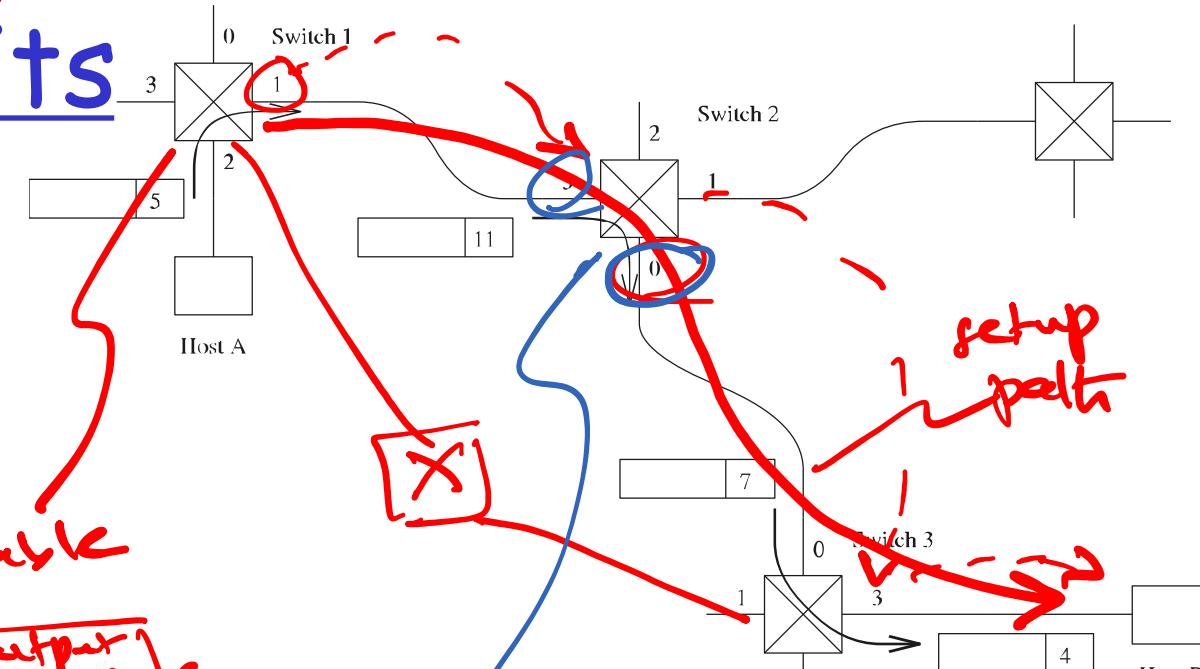
VCI output interface
globally unique
VC #

VCI
data packet

locally unique

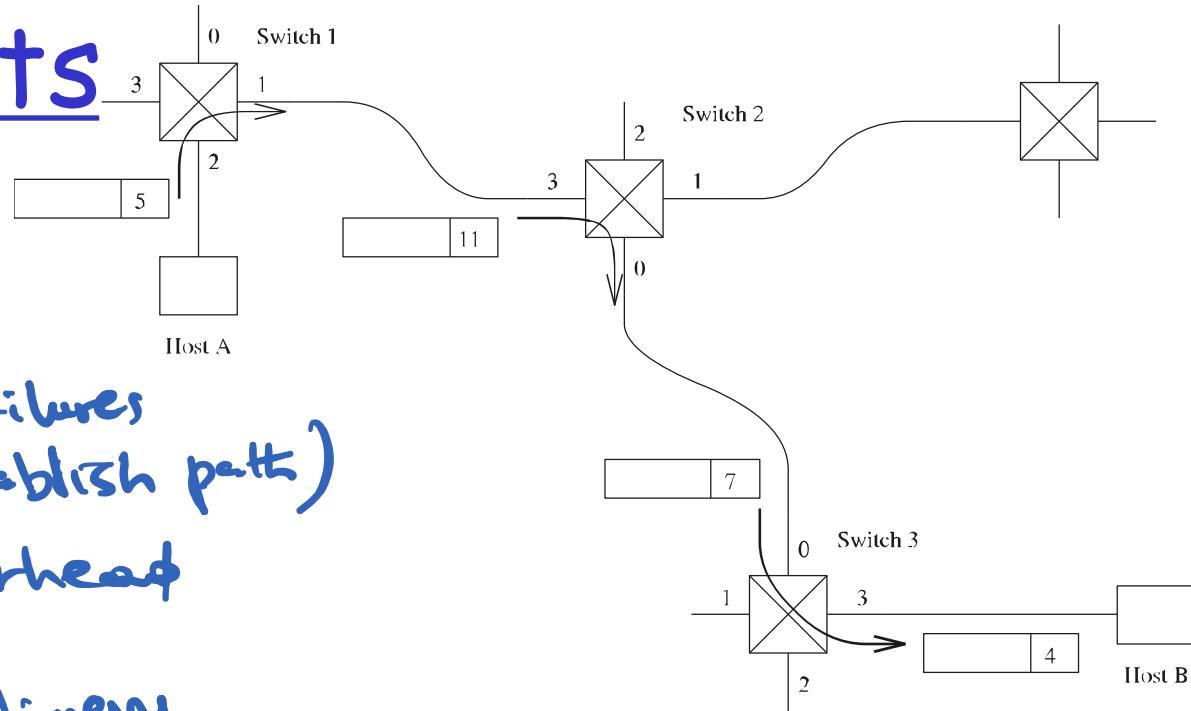
VCI / label swapping

VCI in	iIP interface out	VCI out	old interface
3	y	0	



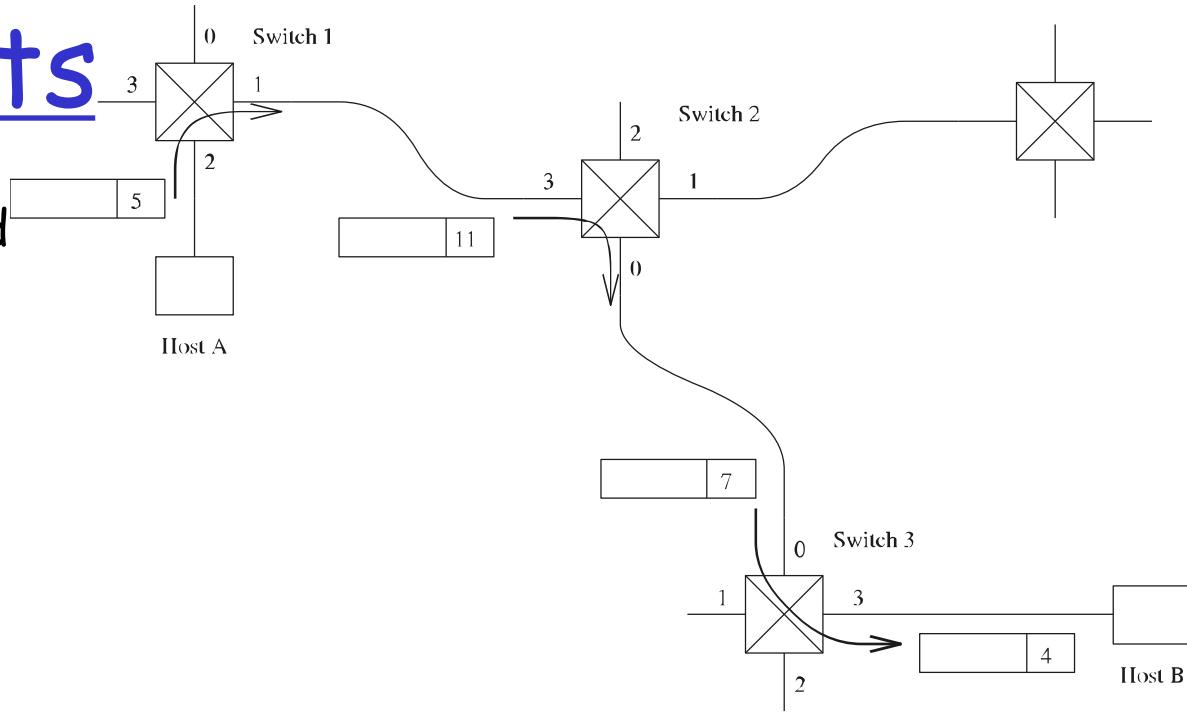
Virtual Circuits

- wait for RTT
 - not resilient to failures
(need to reestablish path)
 - VC tables overhead
-
- + in-order delivery
 - + performance guarantees (QoS)
 - + per data packet overhead smaller



Virtual Circuits

- Explicit connection setup (and tear-down) phase
- Setup packet is routed using next-hop routing
- Subsequent packets follow same route, thus arrive in order
- Analogy: phone call
- Sometimes called **connection-oriented** model (or VC routing)
- Each switch maintains a VC table: {input VCI, input port, output VCI, output port}
- VCI in header is interchanged by nodes as they forward data packet



Virtual Circuits versus Datagrams

Virtual Circuit Model:

- ❑ Typically wait full RTT for connection setup before sending first data packet
- ❑ While the connection request contains the full address of destination, each data packet contains only a small identifier, making the per-packet header overhead small
- ❑ Need per-connection state at switches
- ❑ If a switch or a link in a connection fails, the connection is broken and a new one needs to be established
- ❑ Connection setup provides an opportunity to reserve resources (to guarantee QoS)

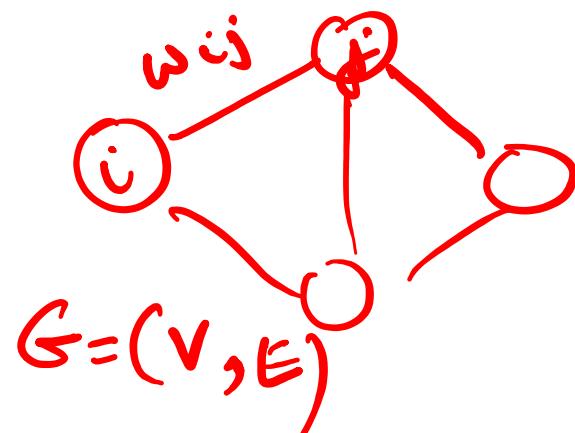
Virtual Circuits versus Datagrams

(cont'd)

Datagram Model:

- There is no round trip time delay waiting for connection setup; a host can send data as soon as it is ready
- Since every packet must carry the full address of the destination, the overhead per packet is higher than for the connection-oriented model
- No connection state in switches
- Since packets are treated independently, it is easier to route around link and node failures
- Source host has no way of knowing if the network is capable of delivering a packet or if the destination host is even up

How do Switches Maintain Routing Tables?



SP problem

Find shortest path P

$$s.t. \sum_{(i,j) \in P} w_{ij} \text{ is min}$$
$$w_{CP} = \sum_{(i,j) \in P}$$

Dest	Next hop
?	?

Distributed routing
Protocols

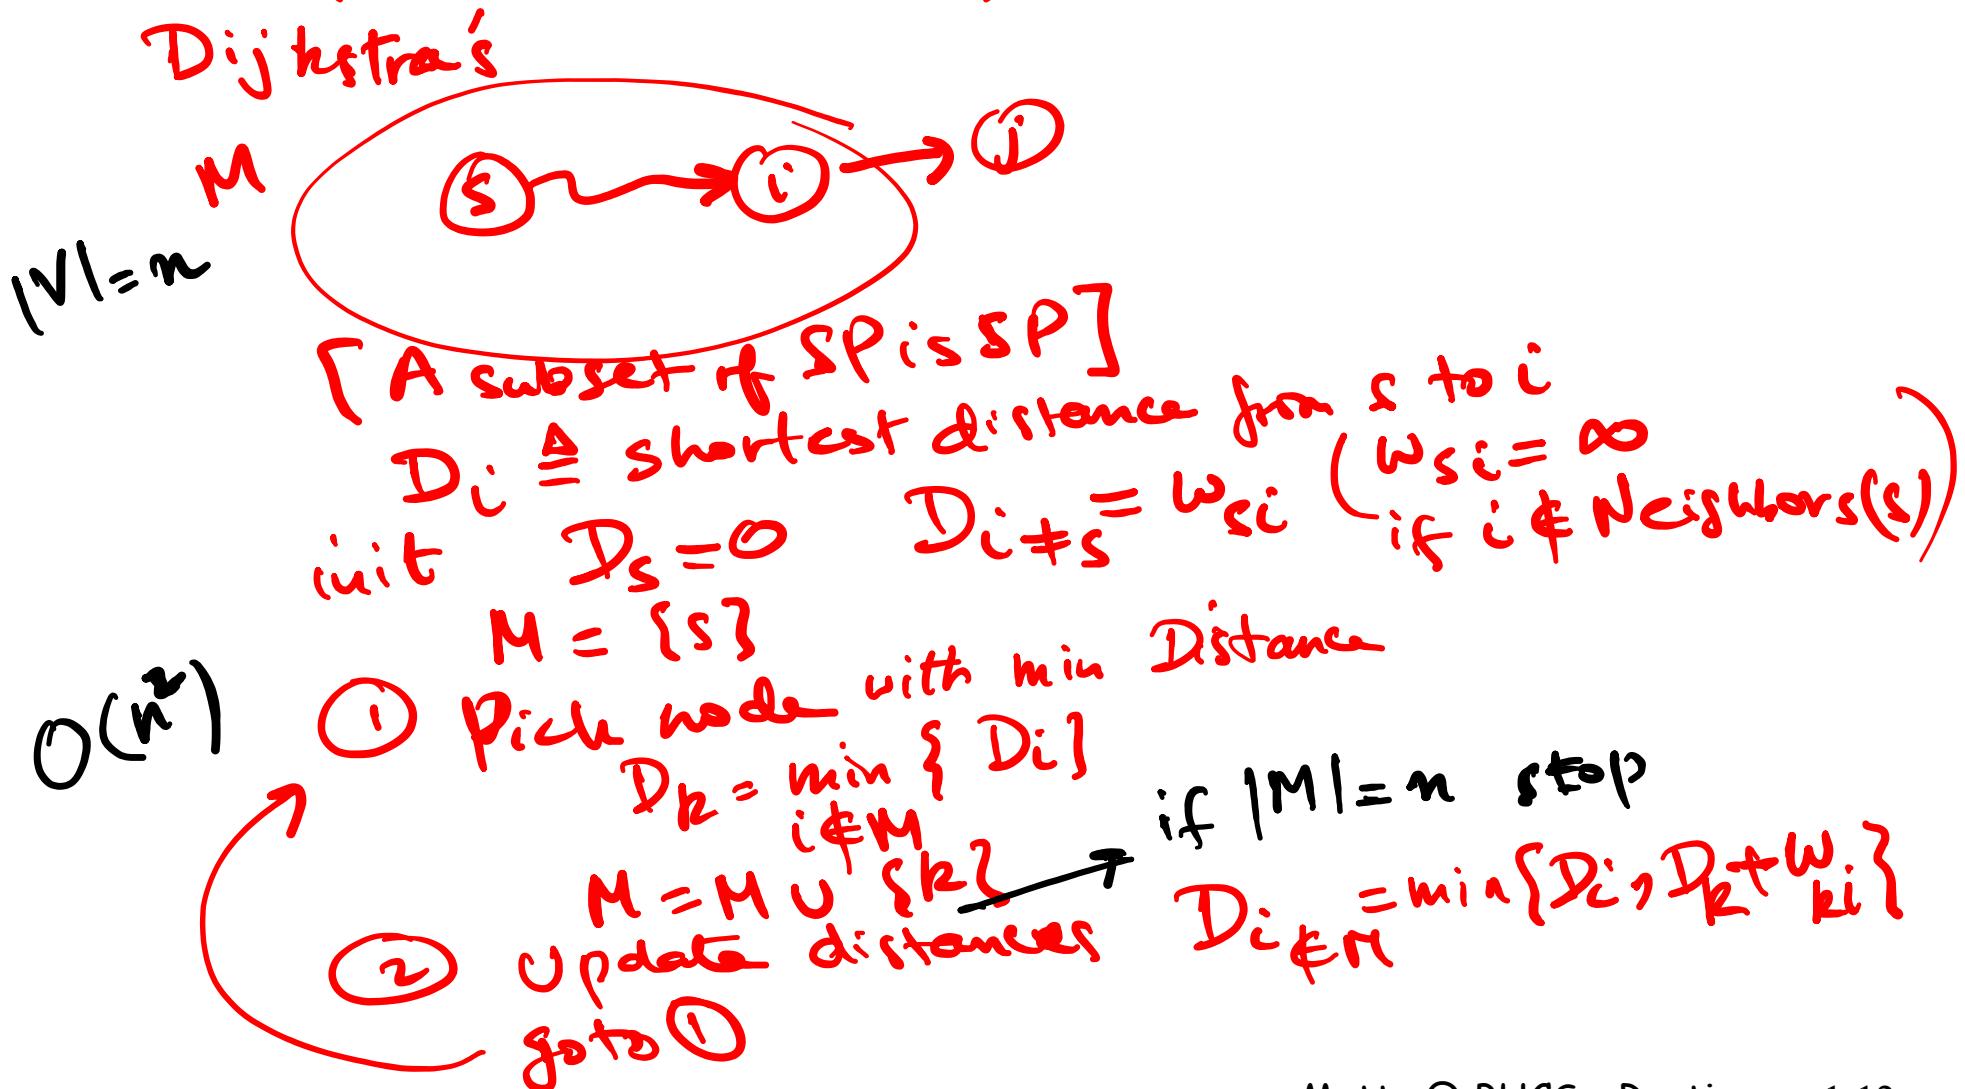
Link-state
(Dijkstra's)

$w_{ij} = \begin{cases} 1 & \text{min-hop routing} \\ \text{delay}_{ij} & \text{min-delay routing} \\ \text{cost}_{ij} & \text{min-cost routing} \end{cases}$

Distance vector
(Bellman-Ford)

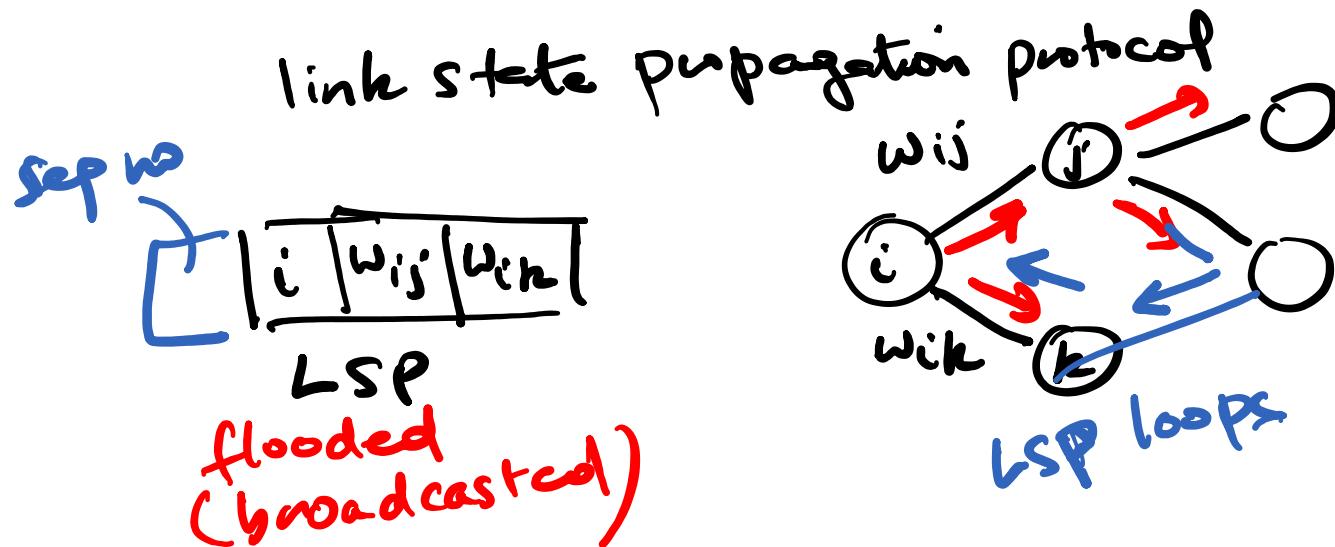
How do Switches Maintain Routing Tables? (cont'd)

Link-state (SPF -- Shortest Path First)



How do Switches Maintain Routing Tables? (cont'd)

Link-state (SPF -- Shortest Path First)

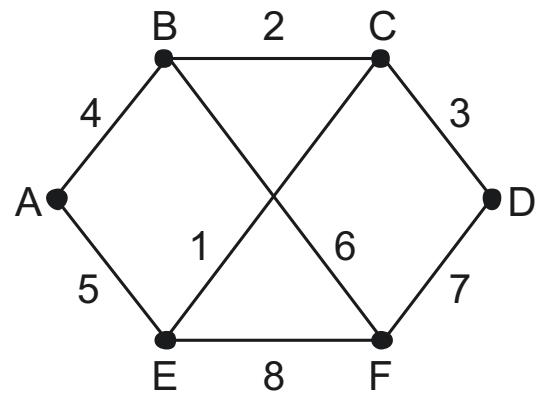


How do Switches Maintain Routing Tables? (cont'd)

Link-state (SPF -- Shortest Path First)

- Each node maintains a complete view of the topology: each node periodically broadcasts to all nodes (not just neighbors) the costs/status of its outgoing links (not entire routing table)
 - Link State Packet (LSP)
 - id of the node that created the LSP
 - cost of link to each directly connected neighbor
 - sequence number (SEQNO)
 - age or time-to-live (TTL) for this packet
- Reliable Flooding
 - store most recent LSP from each node
 - forward LSP to all neighbor nodes but one that sent it
 - generate new LSP periodically; increment SEQNO
 - start SEQNO at 0 when reboot
 - decrement TTL of each stored LSP; discard when TTL=0
- When a link state message arrives and the view changes accordingly, the node recomputes routes by applying "Dijkstra's shortest path algorithm"

LS Propagation Protocol



(a)

	Link	State	Packets	
A	B	C	D	E
Seq.	Seq.	Seq.	Seq.	Seq.
Age	Age	Age	Age	Age
B 4	A 4	B 2	C 3	A 5
E 5	C 2	D 3	F 7	B 6
	F 6	E 1		D 7
				E 8

(b)

Route Calculation in Link State

- Dijkstra's shortest path algorithm
- N denotes set of nodes in the graph
- $l(i,j)$ denotes non-negative cost (weight) for edge (i,j)
- s in N denotes this node
- M denotes the set of nodes incorporated so far
- $C(n)$ denotes cost of the path from s to node n

$$M = \{s\}$$

for each n in $N - \{s\}$

$$C(n) = l(s,n)$$

while ($M \neq N$)

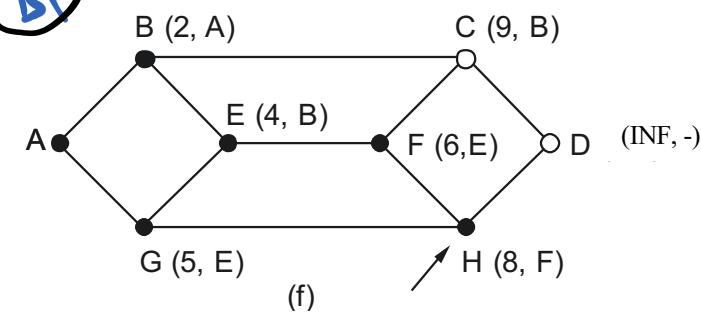
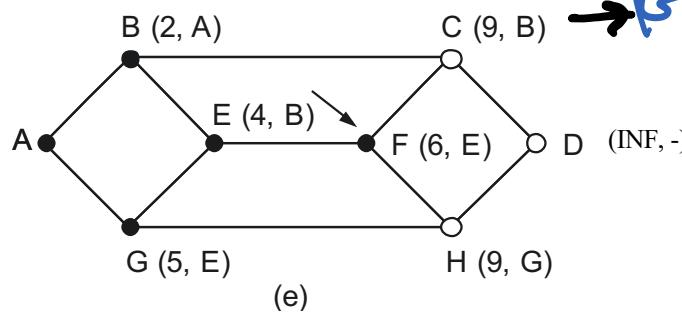
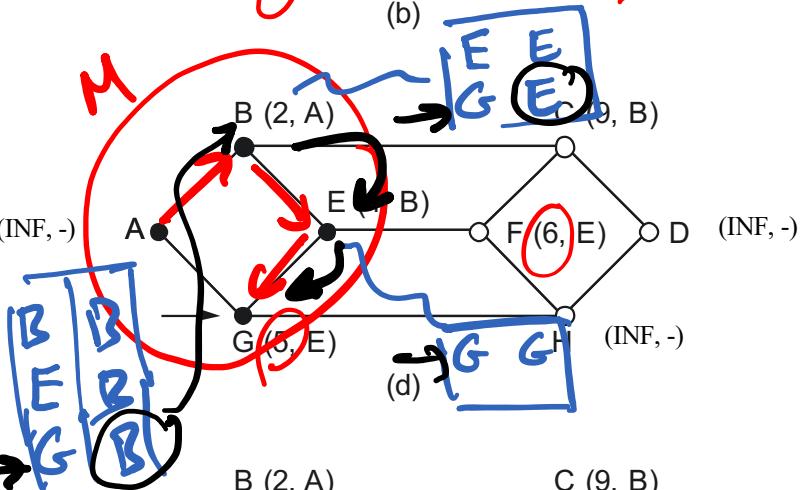
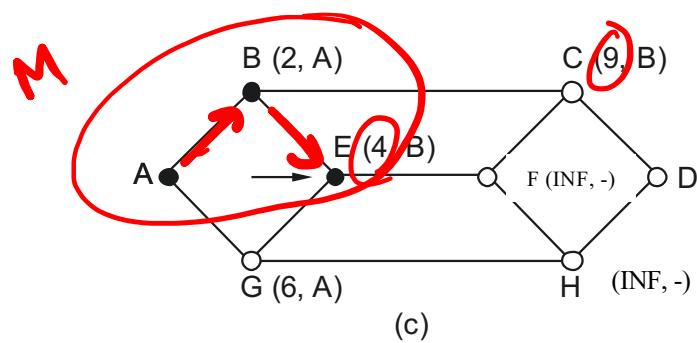
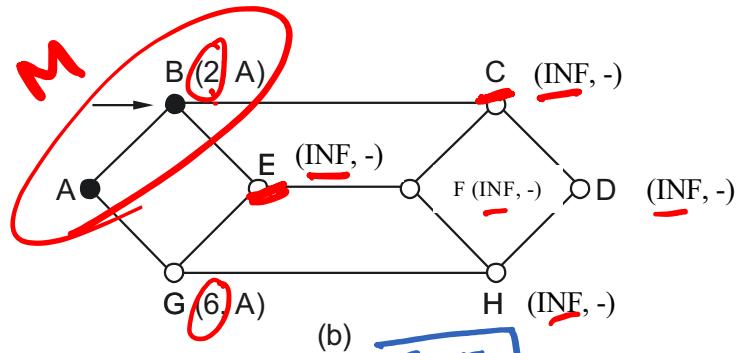
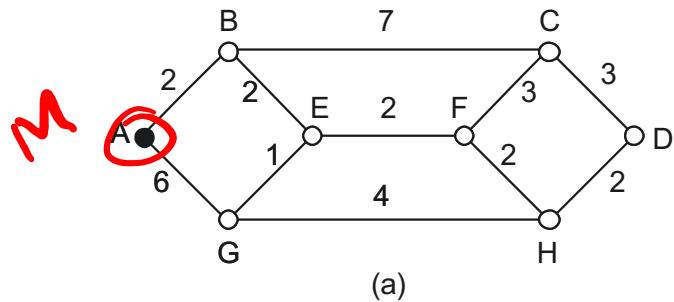
$M = M \cup \{w\}$ such that

$C(w)$ is the minimum for all w in $(N-M)$

for each n in $(N-M)$

$$C(n) = \text{MIN}(C(n), C(w)+l(w,n))$$

Dijkstra's Algorithm in action



How do Switches Maintain Routing Tables?

Distance-vector (Bellman-Ford)

$D_i^h \triangleq$ shortest distance from i to some dest
s.t. length of path $\leq h$ links/hops

init $D_{dest}^h = 0$ $D_{i \neq dest}^0 = \infty$

$|V|=n$ for $h=1, 2, 3, \dots, n-1$
 $D_i^h = \min_{\substack{k \in \\ \text{neighbors}(i)}} \{ D_k^{h-1} + w_{ik} \}$

$O(n^3)$

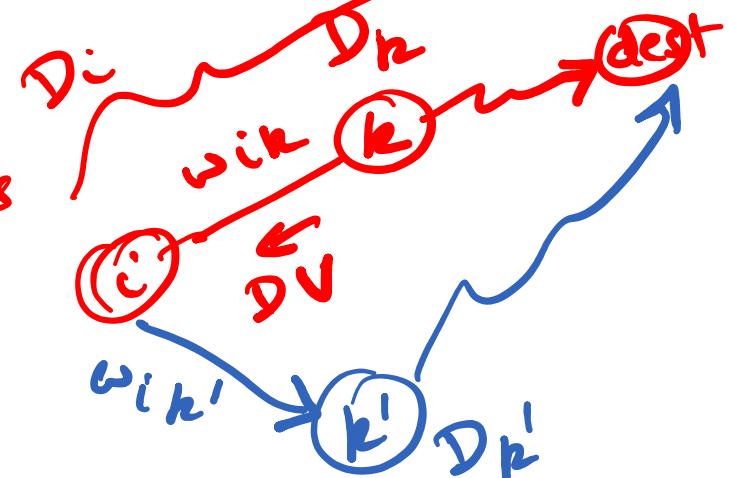
How do Switches Maintain Routing Tables?

Distance-vector (Bellman-Ford)

Periodically, each node sends
a DV = { $(dest, cost)$ }

Upon receipt of DV, node i
updates its routing table if

- 1) unknown dest, take "k" as next hop
- 2) better distance, switch to k as next hop
- 3) k is next hop, update your distance



How do Switches Maintain Routing Tables?

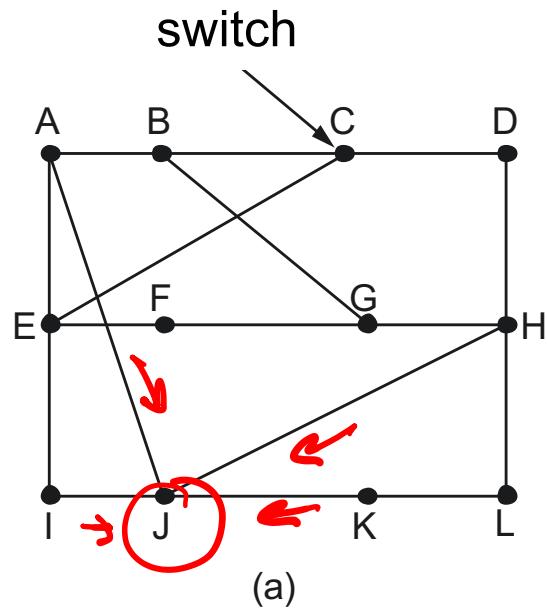
Distance-vector (Bellman-Ford)

- Each node keeps in its routing table its current best distance (cost) to each destination node: (Destination, Cost, NextHop)
- Initially, the routing table of each node contains its neighbor nodes and distances to them
- Each node sends a copy of its routing information {(Destination, Cost)} to each of its neighbor nodes
 - periodically (on the order of several seconds)
 - whenever its table changes (called *triggered update*)
- Upon receiving a message from a neighbor J, a node K updates its routing table as follows:
 - If node J knows a shorter path to some destination, node K updates its corresponding entry with J as the new next-hop
 - If node J reports a distance for an unknown destination, node K adds a new entry for it with J as the next-hop
 - If node K is using node J as its next-hop to some destination node and J reports a change in its distance to it, K updates the corresponding entry

Route Calculation in Distance Vector

```
void mergeRoute (Route *new) {
    int i;
    for (i = 0; i < numRoutes; ++i) {
        if (new->Dest == rt[i].Dest) {
            if (new->Cost + 1 < rt[i].Cost) // min-hop routing
                break; ≡ wij = 1
        else if (new->NextHop == rt[i].NextHop)
            break;
        else
            return;
    }
    rt[i] = *new;
    rt[i].TTL = MAX_TTL;
    ++rt[i].Cost;
    if (i == numRoutes)
        ++numRoutes;
}
```

DV Routing



Distance Table
distance to each dest
via each neighbor
neighbor

dest [48 37 20 30]

DV's

To	A	I	H	K
A	0	24	20	21
B	12	36	31	28
C	25	18	19	36
D	40	27	8	24
E	14	7	30	22
F	23	20	19	40
G	18	31	6	31
H	17	20	0	19
I	21	0	14	22
J	9	11	7	10
K	24	22	22	0
L	29	33	9	9

JA delay is 8
JI delay is 10
JH delay is 12
JK delay is 6

Vectors received from J's four neighbors
48 37 20 30

New estimated delay from J

Line

derived routing table

next hop

New routing table for J

(b)